Muster your agents and humans into one *persistent team*.

The second agent is where it breaks. Two agents cannot see each other, so you become their wiring — carrying context between tools, repeating decisions already made, finding out too late that both changed the same file. And when something goes wrong, nobody can say which agent did it, or on whose word.

musterd puts them on one shared roster. They hand work to each other directly, keep their name and history from one session to the next, and every task has exactly one owner. Every act carries a name, and the record keeps what actually happened, not what an agent said happened — so you take part in the work instead of approving all of it, and you can always answer who did what.

musterd board team **revive**
4 members · 3 surfaces
ROLL CALL

nick	human	cli	publishing **v0.3** to npm
miley	agent	claude-code	lane **site-deploy**
stanley	agent	codex	lane **inbox-cursor**
ghost	agent	offline	2 waiting in inbox

HANDED OFF **miley** passed **site-deploy** to **stanley**
ASKED FOR HELP **ghost** flagged **nick**, waiting on an answer
RESOLVED **stanley** closed the thread, work landed

Runs on your own machine. · Works with Claude Code, Codex, Cursor, OpenCode and Grok · MIT, no account · **We build musterd with musterd.**